

Maximizing LOF As an Objective for Exploration in Reinforcement Learning

Dylan R. Ashley

Outline

- 1. Background

 - 1.1. Local Outlier Factor (LOF)

 - 1.2. Reinforcement Learning

- 2. Motivation

- 3. Basic Idea

- 4. Empirical Analysis

 - 4.1. Experimental Setup

 - 4.1.1. Domain

 - 4.1.2. Algorithms


 - 4.1.3. Performance Metric

 - 4.2. Experimental Results

 - 4.3. Discussion

Local Outlier Factor (LOF)

- Compares the density around a point to the density around each neighbour in the point's k NN
- $LOF \approx 1$ means the point is probably in a cluster
- $LOF \gg 1$ means the point is probably an outlier



Calculating the Local Outlier Factor

$$k\text{-distance}(p) = \max_{q \in k\text{NN}(p)} d(p, q)$$


$$\text{reach-dist}_k(p, q) = \max\{d(p, q), k\text{-distance}(q)\}$$

$$\text{lrd}(p) = \frac{1}{\frac{1}{k} \sum_{q \in k\text{NN}(p)} \text{reach-dist}_k(p, q)}$$

$$\text{LOF}(p) = \frac{\frac{1}{k} \sum_{q \in k\text{NN}(p)} \text{lrd}(q)}{\text{lrd}(p)}$$

Reinforcement Learning

- Models an agent interacting with an environment
- Agent wants to maximize reward signal with later rewards being considered less valuable
- The agent doesn't know about its environment beforehand and has to explore to learn about it



Outline

1. Background
 - 1.1. Local Outlier Factor (LOF)
 - 1.2. Reinforcement Learning
2. Motivation
3. Basic Idea
4. Empirical Analysis
 - 4.1. Experimental Setup
 - 4.1.1. Domain
 - 4.1.2. Algorithms
 - 4.1.3. Performance Metric
 - 4.2. Experimental Results
 - 4.3. Discussion

Motivation for Better Exploration

- Exploration is a major factor in the learning rate for an agent
- Without proper exploration the agent is liable to miss something important
- Naïve approaches often perform very badly



Outline

- 1. Background
 - 1.1. Local Outlier Factor (LOF)
 - 1.2. Reinforcement Learning



- 2. Motivation

- 3. Basic Idea

- 4. Empirical Analysis
 - 4.1. Experimental Setup
 - 4.1.1. Domain
 - 4.1.2. Algorithms
 - 4.1.3. Performance Metric
 - 4.2. Experimental Results
 - 4.3. Discussion

Using the LOF to Guide Exploration

- Reward an RL agent with the LOF of each state or transition it sees and a second agent can learn using the real rewards
- Intuitively this incentivises the agent to find rare states or transitions
- Can combine with random exploration to smooth out the exploration




Outline

1. Background
 - 1.1. Local Outlier Factor (LOF)
 - 1.2. Reinforcement Learning
2. Motivation
3. Basic Idea
4. Empirical Analysis
 - 4.1. Experimental Setup
 - 4.1.1. Domain
 - 4.1.2. Algorithms
 - 4.1.3. Performance Metric
 - 4.2. Experimental Results
 - 4.3. Discussion

Domain

- 5x5 Gridworld with single start and goal state
- Observations drawn from truncated Gaussian centered at current state scaled to cell size leading
- Empirically learnable with simple RL methods




Algorithms

- Q-learning with tile coding as function approximation
- Well known algorithm that learns the value of taking an action based on the state the agent is currently in
- Note that tile coding is actually optimal for this domain

```
Initialize  $Q(s, a)$  arbitrarily
Repeat (for each episode):
  Initialize  $s$ 
  Repeat (for each step of episode):
    Choose  $a$  from  $s$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
    Take action  $a$ , observe  $r, s'$ 
     $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$ 
     $s \leftarrow s'$ ;
  until  $s$  is terminal
```

Algorithms (continued)

- Use a grid for k NN to prevent having to check all data points in k NN search
- Increases overhead and slows early k NN computation in exchange for faster later k NN computation
- No change in worst case time complexity



Performance Metric


$$\sigma = \sqrt{\frac{\sum_{i=1}^{|S|} (V_{S_i} - \bar{V})^2}{|S| - 1}}$$

- Standard deviation of state visitations
- Minimizing means a more even exploration of the world
- Penalizes linear increase in distance to mean superlinearly



Outline

1. Background
 - 1.1. Local Outlier Factor (LOF)
 - 1.2. Reinforcement Learning
2. Motivation
3. Basic Idea
4. Empirical Analysis
 - 4.1. Experimental Setup
 - 4.1.1. Domain
 - 4.1.2. Algorithms
 - 4.1.3. Performance Metric
 - 4.2. Experimental Results
 - 4.3. Discussion

Experimental Results



Experimental Results (continued)



Outline

1. Background
 - 1.1. Local Outlier Factor (LOF)
 - 1.2. Reinforcement Learning
2. Motivation
3. Basic Idea
4. Empirical Analysis
 - 4.1. Experimental Setup
 - 4.1.1. Domain
 - 4.1.2. Algorithms
 - 4.1.3. Performance Metric
 - 4.2. Experimental Results
 - 4.3. Discussion

Discussion

Pros

- Since it's just a reward signal it can be used with any RL algorithm
- Natively handles a continuous space
- Can handle a changing environment by storing only the last n states or transitions

Cons

- Have to store states or transitions
 - If you only store the last n this is viable for the long term
- Doesn't natively handle identical observations
 - Clipping values may be possible without incurring too severe a penalty
- k NN is quite expensive
 - There are cheaper approaches that haven't been considered for this project because of complexity

References

- <https://upload.wikimedia.org/wikipedia/commons/thumb/4/4e/LOF-idea.svg/2000px-LOF-idea.svg.png>
- <http://web.stanford.edu/class/cs234/images/header2.png>
- https://atariage.com/5200/screenshots/s_MontezumasRevenge_2.png
- <http://img.taste.com.au/BcemwldD/taste/2016/11/chocolate-celebration-cake-85607-1.jpeg>
- <https://qph.ec.quoracdn.net/main-qimg-581fc72f6fa620741c9119688fafd5c6>